

Terraform: Infrastructure as Code 2

Aleksandr Usov

Senior System Engineer

HashiCorp Certified: Terraform Associate

Agenda

- Infrastructure as code
- Terraform configuration

Imperative Approach: create

The imperative focuses on how the infrastructure is to be changed

```
$ aws s3api create-bucket --bucket cf08973879519fa5610bc8b6ff6541 --region us-east-2 --create-bucket-configuration LocationConstraint=us-east-2
{
  "Location": "http://cf08973879519fa5610bc8b6ff6541.s3.amazonaws.com/"
}
```

```
$ aws s3api create-bucket --bucket cf08973879519fa5610bc8b6ff6541 --region us-east-2 --create-bucket-configuration LocationConstraint=us-east-2
An error occurred (BucketAlreadyOwnedByYou) when calling the CreateBucket operation: Your previous request to create the named bucket succeeded and you already own it.
$ echo $?
255
```

Imperative Approach: show

```
$ aws s3api list-buckets --region us-east-2
{
  "Buckets": [
    {
      "Name": "cf08973879519fa5610bc8b6ff6541",
      "CreationDate": "2020-04-30T20:05:56.000Z"
    }
  ],
  "Owner": {
    "ID": "905339dfcf0bb1be6066daadd65c3de1799387cf1d6eeb48581860f51ab59c8d"
  }
}
$ aws s3api list-buckets --region us-east-2 | jq .Buckets[0].Name
"cf08973879519fa5610bc8b6ff6541"
```

Imperative Approach: delete

```
$ aws s3api delete-bucket --bucket cf08973879519fa5610bc8b6ff6541 --region us-east-2  
$ aws s3api delete-bucket --bucket cf08973879519fa5610bc8b6ff6541 --region us-east-2
```

```
An error occurred (NoSuchBucket) when calling the DeleteBucket operation: The specified bucket does not exist  
$ echo $?  
255
```

Imperative Approach: example

```
output=$($AWS s3api get-bucket-versioning --bucket $dst | jq '(.Status=="Enabled")')
if [[ $output != true ]]
then
    echo "Enabling versioning for $dst"
    $AWS s3api put-bucket-versioning --bucket $dst --versioning-configuration Status=Enabled
    [[ $? -ne 0 ]] && { echo "Can't enable versioning for $dst"; exit 1; }
fi

$AWS s3api get-bucket-encryption --bucket ${dst} &>/dev/null
if [[ $? -ne 0 ]]
then
    if [[ -z ${key_arn} ]]
    then
        cmk_id=$($AWS kms create-key --origin EXTERNAL --region eu-central-1 | jq '.KeyMetadata.KeyId' | tr -d \")
        [[ $? -ne 0 ]] && { echo "Can't create key"; exit 1; }
        key_arn="arn:aws:kms:eu-central-1:${id}:key/${cmk_id}"
        $AWS kms get-parameters-for-import --key-id ${cmk_id} \
            --wrapping-algorithm RSAES_OAEP_SHA_1 \
            --wrapping-key-spec RSA_2048 --region eu-central-1 >/tmp/get-parameters-for-import
        [[ $? -ne 0 ]] && { echo "Can't download key"; exit 1; }
        openssl enc -d -base64 -A -in PublicKey.b64 -out PublicKey.bin
    fi
fi
```

Declarative Approach

The declarative approach focuses on what the eventual target configuration should be:

```
resource "aws_s3_bucket" "main" {  
    bucket = "cf08973879519fa5610bc8b6ff6541"  
}
```

Declarative Approach: create

```
$ terraform apply -auto-approve
aws_s3_bucket.main: Creating...
aws_s3_bucket.main: Still creating... [10s elapsed]
aws_s3_bucket.main: Creation complete after 12s [id=cf08973879519fa5610bc8b6ff6541]
```

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

Outputs:

```
arn = arn:aws:s3:::cf08973879519fa5610bc8b6ff6541
```

```
$ terraform apply -auto-approve
aws_s3_bucket.main: Refreshing state... [id=cf08973879519fa5610bc8b6ff6541]
```

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.

Outputs:

```
arn = arn:aws:s3:::cf08973879519fa5610bc8b6ff6541
```


Declarative Approach: destroy

```
$ terraform destroy -auto-approve  
aws_s3_bucket.main: Refreshing state... [id=cf08973879519fa5610bc8b6ff6541]  
aws_s3_bucket.main: Destroying... [id=cf08973879519fa5610bc8b6ff6541]  
aws_s3_bucket.main: Destruction complete after 1s
```

```
Destroy complete! Resources: 1 destroyed.  
$ terraform destroy -auto-approve
```

```
Destroy complete! Resources: 0 destroyed.
```

Local Values

A local value assigns a name to an expression, allowing it to be used multiple times within a module without repeating.

Local values can be declared together in a single locals block:

```
locals {  
  company      = "epam"  
  department   = "devops"  
}
```

Data Sources

Data sources allow data to be fetched or computed for use elsewhere in Terraform configuration.

```
data "aws_vpc" "main" {  
  id = "vpc-505d8d3b"  
}  
  
output "vpc_id" {  
  value = data.aws_vpc.main.main_route_table_id  
}
```

```
$ terraform apply:  
...  
Outputs:  
  
vpc_id = rtb-928810f9
```

References to Named Values

The following named values are available:

- `<RESOURCE TYPE>.<NAME>` : an object representing a managed resource
- `var.<NAME>` : the value of the input variable
- `local.<NAME>` : the value of the local value
- `module.<MODULE NAME>.<OUTPUT NAME>` : the value of the specified output value
- `data.<DATA TYPE>.<NAME>` : an object representing a data resource
- `path.module` : the filesystem path of the module where the expression is placed
- `path.root` : the filesystem path of the root module of the configuration
- `path.cwd` : the filesystem path of the current working directory
- `terraform.workspace` : the name of the currently selected workspace

Lifecycle Customizations

The following lifecycle meta-arguments are supported:

- `create_before_destroy (bool)` : a new replacement object is created first, and then the prior object is destroyed only once the replacement is created
- `prevent_destroy (bool)` : will cause Terraform to reject with an error any plan that would destroy the infrastructure object
- `ignore_changes (list of attribute names)` : the given attribute names are considered when planning a create operation, but are ignored when planning an update

Terraform Flow


1. Configuration Loader; Backend; State Manager
2. Create terraform.Context(main object represents all the context that Terraform needs)
3. Graph Builder
4. Graph Walker
5. Vertex Evaluation

Independent Resources

```
resource "aws_iam_group" "testers" {  
  name = "testers"  
  path = "/users/"  
}  
  
resource "aws_s3_bucket" "main" {  
  bucket = "cf08973879519fa5610bc8b6ff6541"  
}
```

Independent Resources

```
$ terraform graph | grep -v meta.count-boundary | dot -Tpng > terraform_independent.png
```




Implicit Dependency

```
resource "aws_s3_bucket" "main" {  
    bucket = "cf08973879519fa5610bc8b6ff6541"  
}  
  
resource "aws_s3_bucket_object" "main" {  
    key      = "someobject"  
    bucket   = aws_s3_bucket.main.id  
    source   = "/etc/debian_version"  
}
```

Implicit Dependency

```
$ terraform graph | grep -v meta.count-boundary | dot -Tpng > terraform_implicit.png
```




Explicit Dependency

```
resource "aws_iam_group" "testers" {  
  name = "testers"  
  path = "/users/"  
}  
  
resource "aws_s3_bucket" "main" {  
  bucket      = "cf08973879519fa5610bc8b6ff6541"  
  depends_on = [aws_s3_bucket.main]  
}
```

Explicit Dependency

```
$ terraform graph | grep -v meta.count-boundary | dot -Tpng > terraform_explicit.png
```



terraform import

The terraform import command is used to import existing resources into Terraform.

Usage: `terraform import [options] ADDRESS ID`

```
$ terraform import aws_s3_bucket.bucket new-bucket

resource "aws_s3_bucket" "bucket" {
  bucket = "new-bucket"
}
```

State in AWS: bootstrap

```
terraform {
  required_version = ">=0.12.24"
  required_providers {
    aws = ">= 2.58.0"
  }
}
provider "aws" {
  region = "eu-central-1"
}
locals {
  tags = {
    source = "terraform"
    env    = var.env
  }
}
resource "aws_s3_bucket" "backend" {
  bucket = "tf-..."
  acl    = "private"

  versioning {
    enabled = true
  }

  tags = local.tags
}
resource "aws_dynamodb_table" "backend_locks" {
  name         = "tf-..."
  billing_mode = "PAY_PER_REQUEST"
  hash_key     = "LockID"

  attribute {
    name = "LockID"
    type = "S"
  }

  tags = local.tags
}
```

State in AWS

```
terraform {  
  backend "s3" {  
    bucket      = "tf-..."  
    encrypt     = true  
    key         = "tf-..."  
    region      = "eu-central-1"  
    dynamodb_table = "tf-..."  
  }  
  required_version = ">=0.12.24"  
  required_providers {  
    aws = ">= 2.58.0"  
  }  
}  
provider "aws" {  
  region = "eu-central-1"  
  assume_role {  
    role_arn = var.admin_role  
  }  
}
```

Backends

Backends are completely optional.

Here are some of the benefits of backends:

- Working in a team
- Keeping sensitive information off disk
- Remote operations

Backend Types:

- **Standard:** State management, functionality covered in State Storage & Locking
- **Enhanced:** Everything in standard plus remote operations.

Enhanced Backends

- local
- remote(Terraform Cloud)

Standard Backends

- artifactory(with no locking)
- azurerm(with state locking)
- consul(with locking)
- etcd(with no locking)
- etcdv3(with locking)
- gcs(with locking)
- http(with optional locking)
- pg(with locking)
- s3(with locking via DynamoDB)
- swift(with no locking)
- terraform enterprise(with no locking)

backend-config

```
backend-config=path
```

This can be either a path to an HCL file with key/value assignments (same format as terraform.tfvars) or a 'key=value' format.

```
$ terraform init -backend-config='bucket=mycompany-tfstate' -backend-config='key=prod.tfstate' \  
-backend-config='region=eu-central-1'
```

S3 Remote State

```
data "terraform_remote_state" "prod" {  
  backend = "s3"  
  config = {  
    bucket = "terraform-state-prod"  
    key     = "terraform.tfstate"  
    region  = "eu-central-1"  
  }  
}
```

The `terraform_remote_state` data source will return all of the root module outputs defined in the referenced remote state.

Modules

Modules help solve the problems:

- **Organize configuration** - Modules make it easier to navigate, understand, and update your configuration by keeping related parts of your configuration together.
- **Encapsulate configuration** - Another benefit of using modules is to encapsulate configuration into distinct logical component.
- **Re-use configuration** - Writing all of your configuration from scratch can be time consuming and error prone.
- **Provide consistency and ensure best practices** - It helps to ensure that best practices are applied across all of your configuration.

Module structure

A typical file structure:

```
$ tree minimal-module/  
.  
├── LICENSE  
├── README.md  
├── main.tf  
├── variables.tf  
└── outputs.tf
```

- `LICENSE` will contain the license under which your module will be distributed.
- `README.md` will contain documentation describing how to use your module.
- `main.tf` will contain the main set of configuration for your module.
- `variables.tf` will contain the variable definitions for your module.
- `outputs.tf` will contain the output definitions for your module.

Modules: nested

```
$ tree complete-module/
```

```
.  
├── README.md  
├── main.tf  
├── variables.tf  
├── outputs.tf  
├── ...  
└── modules/  
    ├── nestedA/  
    │   ├── README.md  
    │   ├── variables.tf  
    │   ├── main.tf  
    │   └── outputs.tf  
    ├── nestedB/  
    └── .../
```

Modules: Local

```
module "consul" {  
  source = "./consul"  
}
```


Modules: Registry

Registry source address: <NAMESPACE>/<NAME>/<PROVIDER>

```
module "consul" {  
  source = "hashicorp/consul/aws"  
  version = "0.1.0"  
}
```

Modules: GitHub

```
module "consul" {  
  source = "github.com/hashicorp/example"  
}
```

Modules: Git

```
module "vpc" {  
    source = "git::https://example.com/vpc.git"  
}  
  
module "storage" {  
    source = "git::ssh://username@example.com/storage.git"  
}
```

terraform-aws-modules

terraform-aws-modules/terraform-aws-vpc/main.tf:

```
#####  
# VPC  
#####  
resource "aws_vpc" "this" {  
  count = var.create_vpc ? 1 : 0  
  
  cidr_block                = var.cidr  
  instance_tenancy          = var.instance_tenancy  
  enable_dns_hostnames      = var.enable_dns_hostnames  
  enable_dns_support        = var.enable_dns_support  
  enable_classiclink         = var.enable_classiclink  
  enable_classiclink_dns_support = var.enable_classiclink_dns_support  
  assign_generated_ipv6_cidr_block = var.enable_ipv6  
  
  tags = merge(  
    {  
      "Name" = format("%s", var.name)  
    },  
    var.tags,  
    var.vpc_tags,  
  )  
}
```

terraform-aws-modules: example

Allow management of default VPC it using Terraform:

main.tf :

```
provider "aws" {  
    region = "eu-west-1"  
}  
  
module "vpc" {  
    source      = "terraform-aws-modules/vpc/aws"  
    version    = "2.33.0"  
  
    create_vpc = false  
  
    manage_default_vpc          = true  
    default_vpc_name            = "default"  
    default_vpc_enable_dns_hostnames = true  
}
```

Registry: Modules

Modules

Modules are self-contained packages of Terraform configurations that are managed as a group.

FILTER BY

Provider



Verified



lb-http
google



Modular Global HTTP Load Balancer for GCE using...

Version 4.0.0 · By GoogleCloudPla...



vpc
aws



Terraform module which creates VPC resources on AWS

Version 2.33.0 · By terraform-aws-...



managed-instance-group
google



Modular Google Compute Engine managed instance group for...

Version 1.1.15 · By GoogleCloudPl...



lb-internal
google



Modular Internal Load Balancer for GCE using forwarding rules.

Version 2.1.0 · By GoogleCloudPla...



nat-gateway
google



Modular NAT Gateway on Google Compute Engine for Terraform.

Version 1.2.3 · By GoogleCloudPla...



loadbalancer
azurerms



Terraform Azure RM Module for Load Balancer

Version 1.2.1 · By Azure

Registry: Requirements

- GitHub. The module must be on GitHub and must be a public repo
- Named `terraform-<PROVIDER>-<NAME>`
- Repository description
- Standard module structure. The module must adhere to the standard module structure
- x.y.z tags for releases

terraform: state

The terraform state command is used for advanced state management.

Usage: `terraform state <subcommand> [options] [args]`

- `terraform state list` command is used to list resources within a Terraform state
- `terraform state mv` command is used to move items in a Terraform state
- `terraform state pull` command is used to manually download and output the state from remote state
- ...

terraform: state

- `terraform state push` command is used to manually upload a local state file to remote state
- `terraform state rm` command is used to remove items from the Terraform state
- `terraform state show` command is used to show the attributes of a single resource in the Terraform state

Terraformer

A CLI tool that generates tf/json and tfstate files based on existing infrastructure (reverse Terraform)

```
$ terraformer import aws --resources=vpc,subnet
2020/05/02 20:56:49 aws importing default region
2020/05/02 20:56:49 aws importing... vpc
2020/05/02 20:56:56 Refreshing state... aws_vpc.tferr--vpc-002D-505d8d3b
2020/05/02 20:57:05 aws importing... subnet
2020/05/02 20:57:12 Refreshing state... aws_subnet.tferr--subnet-002D-d1bc47ba
2020/05/02 20:57:12 Refreshing state... aws_subnet.tferr--subnet-002D-0e487974
2020/05/02 20:57:12 Refreshing state... aws_subnet.tferr--subnet-002D-46a0390a
2020/05/02 20:57:19 aws Connecting....
2020/05/02 20:57:19 aws save vpc
2020/05/02 20:57:19 aws save tfstate for vpc
2020/05/02 20:57:19 aws save subnet
2020/05/02 20:57:19 aws save tfstate for subnet
```

Terraformer

```
$ cat generated/aws/subnet/variables.tf
data "terraform_remote_state" "vpc" {
  backend = "local"

  config = {
    path = "../../../generated/aws/vpc/terraform.tfstate"
  }
}
$ cat generated/aws/subnet/subnet.tf
resource "aws_subnet" "tfer--subnet-002D-0e487974" {
  assign_ipv6_address_on_creation = "false"
  cidr_block                      = "172.31.16.0/20"
  map_public_ip_on_launch        = "true"
  vpc_id                         = "${data.terraform_remote_state.vpc.outputs.aws_vpc_tfer--vpc-002D-505d8d3b_id}"
}
```

Links

[Terraform Documentation](#)

End